Synchronic Semantic Variation Across Groups

Project Report

Johnny Meng and Yimo Ning
{johnnymeng, yimoning}@cs.toronto.edu

Abstract

Semantic variation arises not only over time, but also across social groups and domains. In this work, we study synchronic semantic variation by using diachronic semantic change methods in a group-based setting. We construct four aligned corpora from different domains: arXiv, Yelp, Ciao, and Reddit, and restrict them to the same time period to enable synchronic analysis.

We train word embeddings using word2vec, PPMI, and SVD, and align embedding spaces via orthogonal procrustes. Semantic variation is derived using cosine distance, and consistency across methods is evaluated using Spearman correlation. We further apply the idea from ConShift to identify stable anchor words and find words with variation.

Results show that PPMI and SVD produce consistently correlated signals, while word2vec captures different patterns of semantic variation.

The code for this project is available at https://github.com/BullDF/synchronic-semantic-variation.

1 Introduction

Recent work shows that semantic differences can arise synchronically across social groups and discourse contexts, even when time is constant. Words may be framed and interpreted differently in political and media discourse, across polarized communities, and within online communities of practice. These observations suggest that semantic variation not only happens via time change but is also shaped by social, cultural, and domain-specific factors.

There is no framework for systematically measuring and comparing such variation in a purely synchronic setting for different social groups. Most existing frameworks for systematically detecting semantic change are diachronic, where time is the main axis, and it remains unclear how well these methods generalize when time is fixed and variation arises across groups instead.

In this work, we use diachronic semantic change detection methods in a synchronic, group-based setting. Specifically, we replace time-based analysis with corpora representing different domains and social groups at the same time frame, and apply multiple embedding-based techniques to quantify cross-group semantic variation. We construct comparable datasets from arXiv, Yelp, Ciao, and Reddit, all selected from the same time period, and analyze semantic differences using aligned embedding spaces and cosine distance.

Our experiments show that count-based methods such as PPMI and SVD produce consistently correlated signals of semantic variation, while neural embeddings such as word2vec capture different patterns across corpora. Overall, this work provides a systematic empirical study of synchronic semantic variation.

2 Related Work

2.1 Diachronic Semantic Change

Most studies on detecting semantic change have focused on diachronic settings, where word meanings are compared across different time periods and time is the primary axis of variation. Earlier research used statistical methods to identify changes in word usage distributions, specifically to detect semantic change by tracking shifts in word frequency and co-occurrence patterns over time. (Kulkarni et al., 2015) Some work has used both distributional and embedding-based representa-

tions to quantify semantic change over time. Hamilton et al. (2016b) showed that diachronic word embeddings capture systematic patterns of semantic change, including relationships between frequency, polysemy, and rates of change. However, later work shows that some observed patterns may be influenced by properties of the models rather than actual semantic shifts. (Dubossarsky et al., 2017).

More recent methods incorporate contextualized representations to capture sense-specific semantic change. For example, dynamic contextual embeddings extend traditional models by modeling both temporal and contextual variation, which allows for more detailed analysis of semantic change (Hofmann et al., 2021a). While these approaches work well for modeling temporal change, they assume gradual variation over time and may not generalize to settings where variation occurs across groups within the same time period.

2.2 Synchronic Semantic Variation Across Groups

In addition to temporal change, studies have shown that semantic variation can occur across social groups, domains, and discourse contexts. For instance, Azarbonyad et al. (2017) showed that word meanings shift across political and media posts, while KhudaBukhsh et al. (2021) examined semantic polarization by analyzing semantic differences between ideological communities. Similarly, Del Tredici and Fernández (2017) studied semantic variation in online communities, showing how community-specific practices influence word usage.

Other work emphasizes the role of social and cultural structures in shaping meaning. Noble et al. (2021) showed that semantic variation correlates with social network organization, while cross-cultural studies demonstrate that word meanings vary systematically across cultures while maintaining shared structure (Jackson et al., 2019; Thompson et al., 2020). These findings suggest that semantic representations are influenced by both universal patterns and group-specific factors.

Recent work has also explored structured approaches to analyzing semantic variation in a synchronic setting. For example, Cassotti et al. (2024) leverages synchronic definitions and semantic relations to classify types of semantic change, highlighting the importance of systematic comparison within a fixed time frame.

2.3 Structured and Domain-Specific Semantic Variation

A related line of work examines semantic variation along specific dimensions such as sentiment and semantic orientation. Methods based on distributional association and PPMI have been used to detect changes in word polarity over time (Cook and Stevenson, 2010) and to build domain-specific sentiment lexicons from unlabeled corpora (Hamilton et al., 2016a). These approaches provide interpretable mechanisms for capturing structured variation, although they are typically applied in diachronic or domain-shift settings.

In addition, prior studies have shown that semantic representations encode social biases and their evolution over time. For example, Garg et al. (2018) and Charlesworth et al. (2022) show that word embeddings capture long-term changes in how social groups are represented, and also show how meaning varies across social contexts.

2.4 Anchor Word Identification and Seed-Based Methods

Several prior approaches rely on seed words or manually defined anchors to support semantic analysis. Earlier work on sentiment orientation detection uses predefined seed words to propagate sentiment and identify changes in polarity (Cook and Stevenson, 2010). Similarly, domain-specific sentiment lexicon induction methods expand from a small set of seed words to capture semantic variation across domains (Hamilton et al., 2016a). These methods are effective but rely on manually selected seeds, which can restrict generalization across different tasks and domains.

In embedding-based approaches, alignment methods often rely on implicitly stable words as anchors. For example, diachronic embedding alignment assumes that a subset of vocabulary remains semantically stable across corpora, enabling comparison through techniques such as Orthogonal Procrustes (Hamilton et al., 2016b). However, such anchors are not explicitly identified and may introduce

noise when applied to heterogeneous corpora.

More recent work explores structured methods for comparing semantic variation without explicit seed words, such as leveraging definitions and semantic relations to characterize changes (Cassotti et al., 2024). Building on these ideas, our work employs the ConShift algorithm (Arrington et al., 2025), which iteratively identifies anchor words in a data-driven manner and refines alignment across corpora. This approach avoids reliance on predefined seeds while providing a principled mechanism for distinguishing stable and variable words.

3 Methodology

3.1 Data

In the study conducted by Hofmann et al. (2021b), the researchers extended BERT's contextualized embeddings by incorporating the time and social space dimensions, rendering word embeddings *dynamic*. To evaluate the model's performance on different social groups, the authors used text datasets from different social contexts. These datasets are publicly available online, making them suitable for the current study. Below is a description of each corpus we used in the project:

3.1.1 arXiv

arXiv is currently one of the largest repositories that hosts scientific research papers in various subjects, with a major focus on computer science. A dataset containing the metadata of papers stored on arXiv is publicly available here on Kaggle. To keep the dataset focused, we filtered the papers by domains and extracted the abstracts of only papers in computer science (category starting with `cs.`). By training embedding models on arXiv abstracts, we aimed to capture semantic variation for words specifically used in an academic, scientific setting.

3.1.2 Yelp

Yelp is a restaurant review platform where users can post reviews for businesses around the world. A comprehensive dataset containing reviews, photos, check-in information, and other attributes of businesses partnered with Yelp is available here on the official website of Yelp. From the dataset, we extracted all user review texts for analyzing semantic variation in a business domain.

3.1.3 Ciao

Ciao is an online review site, similar to Yelp, but focused on product reviews rather than restaurant reviews. A public dataset containing reviews on Ciao during 2000-2011 is made available here on the website of a professor at Michigan State University. One characteristic of this Ciao dataset is that the reviews are relatively long compared to other corpora used in this project: the mean length of the reviews is 778 tokens, and the maximum length is 14k tokens.

3.1.4 Reddit

Lastly, Reddit is an online forum where people discuss various topics, classified under different communities called "subreddits". Datasets containing a large amount of Reddit posts and comments during 2005-2025 are available here on the site of Academic Torrents, with a general dataset available here. Since the datasets are very large, we did not filter the data by subreddits like what we did for arXiv, and instead, we pulled all the Reddit posts and excluded the comments.

Table 1 shows an overview of the sizes of the corpora described above, with arXiv being the smallest and Reddit being substantially larger. This table suggests a severe data imbalance in our datasets that may hinder performance and a fair comparison between models in our results. Additionally, since we are examining synchronic semantic variation, the datasets must be gathered from the same time period so that the comparison is meaningful. However, the Ciao dataset is only available during 2000-2011, while the majority of arXiv papers are from 2010 onwards. Therefore, we had to make the compromise to restrict all 4 datasets to 2010-2011, and we acknowledge this as a data limitation.

3.2 Embedding Algorithms

In this paper, we employed 3 methods to train word embeddings: word2vec, positive pointwise mutual information (PPMI), and singular value decomposition (SVD). For each corpus, all 3 embedding methods were applied, leading to 3 different embedding spaces in which we

Corpus	Domain	Documents
arXiv CS	Academic (CS)	16,711
Yelp	Restaurant reviews	369,400
Ciao	Product reviews	48,018
Reddit	Online communities	16M

Table 1: Corpora used for cross-group comparison

calculated cosine distances to measure semantic variation. This design allowed us to more confidently identify semantic variation when multiple methods agreed. This section briefly discusses the theories behind each of the embedding algorithms.

3.2.1 word2vec

word2vec is a classical word embedding training mechanism introduced by Mikolov et al. (2013). There are two training methods for word2vec: continuous bag of words (CBOW) and skip-gram with negative sampling (SGNS). We used SGNS in this project, where given a target word w_i, the model was asked to predict the context words c_j surrounding w_i. The training maximized the probability of true context words while minimizing those of sampled "negative" words unlikely to appear in the context of the target word. At the end, each word w_i is represented by a word vector $\mathbf{w}_i^{\text{SGNS}}$ and a context vector $\mathbf{c}_j^{\text{SGNS}}$ that satisfy

$$\mathbb{P}(c_j|w_i) \propto \exp(\mathbf{w}_i^{\text{SGNS}} \cdot \mathbf{c}_j^{\text{SGNS}}).$$

In this project, we used a context window size of 4, following Hamilton et al. (2016b).

3.2.2 PPMI

PPMI is a count-based method (Levy et al., 2015). Given word w_i and a fixed context window size, we compute the positive pointwise mutual information of w_i with the set of context words c_j appearing within the context window size using the following formula:

$$\mathbf{M}_{i,j}^{\text{PPMI}} = \max\left\{\log \frac{\mathbb{P}(w_i, c_j)}{\mathbb{P}(w_i)\mathbb{P}(c_j)} - \alpha, 0\right\},$$

where $\alpha > 0$ is a shift parameter (Hamilton et al., 2016b). In this study, we used a context window size of 4 and $\alpha = \log 5$, following Hamilton et al. (2016b). Generally, the PPMI matrix is sparse because only certain words in the vocabulary co-occur with the target word w_i, and this might lead to problems in computing cosine distances. More on this will be discussed in *Section 4* and *Section 5.2*.

3.2.3 SVD

SVD is a technique in linear algebra that decomposes a matrix into 3 components, mathematically expressed as

$$\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top.$$

In the equation, the columns of matrix $\mathbf{U}$ are orthonormal vectors representing the principal directions of the space, and each row of $\mathbf{U}$ corresponds to a word. Matrix $\mathbf{\Sigma}$ is a diagonal matrix where each value on the diagonal represents the importance of the corresponding principal direction. To train word embeddings using SVD, we applied SVD on the PPMI matrix described in the previous section (however, with $\alpha = 0$ following Hamilton et al. (2016b)) and obtained the embedding for word w_i as

$$\mathbf{w}_i^{\text{SVD}} = (\mathbf{U}\mathbf{\Sigma}^\gamma)_i,$$

where γ is a weighting parameter. We discarded $\mathbf{V}^\top$ and used only $\mathbf{U}$ and $\mathbf{\Sigma}$ to form the embeddings. We used $\gamma = 0.5$ for the SVD training (Hamilton et al., 2016b).

3.3 ConShift

We apply `ConShift` (Arrington et al., 2025) as an iterative procedure to identify anchor (low-shift) and shifted (high-shift) words across corpus pairs. Using the aligned embeddings described above, we compute cosine distance as a shift score for each shared word.

At each iteration, low-shift words are selected as anchors and used to re-align the embedding spaces. Shift scores are then recomputed, and the process repeats until convergence. The final low-shift set is treated as anchor words, while high-shift words are taken as candidates for semantic variation.

In addition, to identify words that are consistently stable across all corpora, we aggregate pairwise shift scores by computing the mean shift of each word across all corpus pairs and rank words by lowest mean shift. This provides a simple estimate of global semantic stability.

3.4 Experimental Design

With the selected datasets and embedding algorithms, we started the experiments by training the embeddings using each embedding algorithm on each corpus, resulting in 12 embedding models. However, the embedding algorithms described above have a known problem, primarily associated with word2vec and SVD. Particularly, depending on the model initialization, we might obtain word2vec embeddings with different orientations in the embedding space. The singular value decomposition of a matrix is also not unique if a singular value has a multiplicity greater than 1. Therefore, to make the embeddings comparable across corpora, we had to use the Procrustes alignment algorithm to align the word2vec and SVD embeddings of different corpora. This alignment step assumes that a subset of words remains semantically stable across corpora, allowing the embedding spaces to be mapped into a shared coordinate system. Without alignment, differences in embedding orientation would make direct comparison meaningless, as cosine distances would reflect arbitrary rotations rather than true semantic variation. Following Hamilton et al. (2016b), let $\mathbf{W}_i$, $\mathbf{W}_j$ be the embedding matrices for corpora i and j. We first L2-normalize the two embedding matrices to get $\mathbf{W}_i^{\text{norm}}$ and $\mathbf{W}_j^{\text{norm}}$, so that all word vectors have unit length. Then, we found the best rotational matrix by optimizing

$$\mathbf{R} = \arg \min_{\mathbf{Q}^\top \mathbf{Q} = \mathbf{I}} ||\mathbf{W}_i^{\text{norm}} \mathbf{Q} - \mathbf{W}_j^{\text{norm}}||_F.$$

Note that PPMI does not need alignment because it is fully count-based. After the Procrustes transformation, we could then meaningfully compare vectors trained on different corpora of the same words by computing the cosine distance between them as

$$\text{cos-dist}(\mathbf{w}_i, \mathbf{w}_j) = 1 - \frac{\mathbf{w}_i \cdot \mathbf{w}_j}{||\mathbf{w}_i||_2 ||\mathbf{w}_j||_2}.$$

With the cosine distances computed for each embedding method between each pair of corpora, we examined the Spearman correlation between cosine distances computed from different embedding methods. The significance is that, if the correlation is positive and strong, then we have strong evidence that these embedding algorithms collectively detect meaningful semantic variation across these social groups. On the other hand, if the correlation is weak or negative, then there might be different signals that these algorithms are detecting.

Pair	w2v vs PPMI	w2v vs SVD	PPMI vs SVD
arXiv vs Yelp	−0.07	−0.02	0.44
arXiv vs Ciao	−0.03	0.04	0.51
arXiv vs Reddit	**0.29**	**0.50**	0.50
Yelp vs Ciao	−0.15	−0.23	**0.61**
Yelp vs Reddit	0.11	0.08	0.60
Ciao vs Reddit	0.14	−0.02	0.54

Table 2: Spearman ρ on common vocabulary per pair

Lastly, we employed an iterative algorithm called `ConShift` to find anchor words from the computed cosine distances (Arrington et al., 2025). This algorithm finds both a set of stable words and a set of words that it determines to have shifted in meaning, which allows us to conduct qualitative analysis of semantic variation. More on this algorithm is described in the next section.

4 Results & Discussion

4.1 Spearman Correlation

Table 2 shows the Spearman correlations between each pair of the embedding algorithms on each pair of corpora. From the rightmost column of the table, we see that the Spearman correlations for PPMI and SVD are all positive and moderately strong. This outcome is expected, as SVD is derived from PPMI as mentioned above, and so it is highly probable that they capture similar signals. It is also worth mentioning that the pair of arXiv and Reddit achieved the highest Spearman correlations on both the word2vec and PPMI pair and the word2vec and SVD pair. We attribute this to arXiv and Reddit falling under similar domains, as the Reddit dataset used in this study contains computer science-related subreddits. Other than these, the Spearman correlations for other pairs of corpora are near 0, indicating that the embedding methods were capturing different information on these datasets, and thus any conclusion of semantic variation in these pairs of corpora should be taken with caution.

One explanation for the near-zero Spearman correlations involving PPMI might be related to the sparse nature of the PPMI embedding matrix. Looking at the PPMI formula pre-

Pair	% dist $\in [0.95, 1.05]$	Vocab size
arXiv vs Yelp	70.6%	8,810
arXiv vs Ciao	64.6%	9,540
arXiv vs Reddit	44.7%	12,640
Yelp vs Ciao	47.5%	33,854
Yelp vs Reddit	41.0%	48,508
Ciao vs Reddit	33.9%	55,164

Table 3: Percentage of words with PPMI cosine distance $\in [0.95, 1.05]$

sented previously, the PPMI matrix is mostly count-based, where the embeddings are computed based on co-occurrences of words. In this sense, if the context words appearing near a target word are drastically different in different corpora, then the dot product of these vectors would be near 0, leading to cosine distances of near 1 for most words in the vocabulary. To quantify this, *Table 3* reports the percentage of words with PPMI cosine distance in $[0.95, 1.05]$ for each pair, indicating how many words have cosine distances clustered near 1. Notably, the highest percentage of 70.6% occurs between arXiv and Yelp, and even the lowest percentage is 33.9% between Ciao and Reddit, despite having the largest common vocabulary. These numbers suggest that PPMI embeddings are unreliable for analyzing synchronic semantic variation across distant domains. More on this will be mentioned in *Section 5.2*.

4.2 `ConShift` Analysis

We apply ConShift to identify stable (low shift) and variable (high shift) words across corpus pairs, and further aggregate shift scores by computing the mean shift of each word across all pairs to obtain globally stable words. Across methods, stable words tend to be general and context independent. For example, *difficult*, *larger*, and *students* exhibit consistently low shift scores across domains. PPMI and SVD produce similar sets of stable words, which is expected since SVD is derived from the PPMI matrix and captures similar distributional information. In contrast, word2vec produces different stable word sets, often including corpus-specific terms such as proper nouns, reflecting differences in training objectives.

High shift words provide clear evidence of semantic variation. For instance, *mays* refers to a location or restaurant name in Yelp, while in Ciao it appears in ingredient lists (e.g., *zea mays*, referring to corn). Similarly, *twitch* in Yelp describes a physical reaction, whereas in Reddit it refers to the online streaming platform. These examples show how word meanings vary across domains, supporting our overall finding that semantic variation depends on both data and modeling assumptions.

5 Conclusion

5.1 Summary

To summarize, this paper aimed to apply methods for detecting diachronic semantic variation to a synchronic setting and find semantic variation there. We adapted the procedure carried out by Hamilton et al. (2016b) and applied it to the corpora used by Hofmann et al. (2021b), consisting of arXiv CS paper abstracts, Ciao, Yelp, and Reddit during the 2010-2011 period. Our results showed that cosine distances computed using PPMI and SVD were highly correlated, which is expected as SVD is based on the embedding matrix computed using PPMI. The correlations between word2vec and other methods were minimal on most pairs of corpora besides arXiv and Reddit, suggesting that word2vec, as a neural model, captures different semantic variation signals than count-based methods such as PPMI and SVD. We further noted that the PPMI embedding method has a significant limitation that collapses a large portion of cosine distances to 1 for distant domains, as shown in *Table 3*. Lastly, we conducted qualitative analysis via the `ConShift` algorithm to find anchor words and words that had undergone semantic variation between corpora (Arrington et al., 2025). From the results, we found words such as *difficult*, *larger*, and *issues* that have relatively stable meanings across corpora, while words including *excited* and *valued* demonstrated successful semantic variation detection, demonstrating the usefulness of our embedding methods and justifying our experimental procedure.

5.2 Limitations & Future Work

As mentioned above, PPMI poses a significant limitation in our analysis. The reason that it

worked in the diachronic setting of Hamilton et al. (2016b) is that time is a continuous dimension. Since semantic shift along the time axis is gradual, the possible context words around a target word are also gradually changing, and thus, we would not encounter similar cases as in our study, where the context words for different corpora are drastically different. Considering this, future researchers studying PPMI in a synchronic setting can consider including corpora that are more granular, such as similar subreddits. This is because the corpora in our study are too distinct from each other, and by including more granular social groups, one could exploit the "continuous" nature of social groups to attempt to make PPMI more useful.

Another issue with the current study is the difficulty of distinguishing semantic variation from polysemy. For example, the word *transformer* has completely different meanings in a movie context versus a computer science context due to polysemy, but our methods would classify it as semantic variation since this word would have different embedding vectors in different corpora. To address this, future studies might consider using more sophisticated embedding methods, such as BERT, so that the models can be aware of the context in order to make a decision.

6 Acknowledgements

We are extremely grateful to Professor Yang Xu and TA Kieran Henderson for giving valuable feedback and conceptual guidance in computational semantic change throughout the process of this research.

References

Clare Arrington, Maurício Gruppi, and Sibel Adali. 2025. Conshift: Sense-based language variation analysis using flexible alignment. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 167–181.

Hosein Azarbonyad, Mostafa Dehghani, Kaspar Beelen, Alexandra Arkut, Maarten Marx, and Jaap Kamps. 2017. Words are malleable: Computing semantic shifts in political and media discourse. In *Proceedings of the 2017 ACM on Conference on Information and Knowledge Management*, pages 1509–1518.

Pierluigi Cassotti, Stefano De Pascale, and Nina Tahmasebi. 2024. Using synchronic definitions and semantic relations to classify semantic change types. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4539–4553.

Tessa ES Charlesworth, Aylin Caliskan, and Mahzarin R Banaji. 2022. Historical representations of social groups across 200 years of word embeddings from google books. *Proceedings of the National Academy of Sciences*, 119(28):e2121798119.

Paul Cook and Suzanne Stevenson. 2010. Automatically identifying changes in the semantic orientation of words. In *LREC*.

Marco Del Tredici and Raquel Fernández. 2017. Semantic variation in online communities of practice. In *Proceedings of the 12th International Conference on Computational Semantics (IWCS)–Long papers*.

Haim Dubossarsky, Daphna Weinshall, and Eitan Grossman. 2017. Outta control: Laws of semantic change and inherent biases in word representation models. In *Proceedings of the 2017 conference on empirical methods in natural language processing*, pages 1136–1145.

Nikhil Garg, Londa Schiebinger, Dan Jurafsky, and James Zou. 2018. Word embeddings quantify 100 years of gender and ethnic stereotypes. *Proceedings of the National Academy of Sciences*, 115(16):E3635–E3644.

William L Hamilton, Kevin Clark, Jure Leskovec, and Dan Jurafsky. 2016a. Inducing domain-specific sentiment lexicons from unlabeled corpora. In *Proceedings of the conference on empirical methods in natural language processing. conference on empirical methods in natural language processing*, volume 2016, page 595.

William L Hamilton, Jure Leskovec, and Dan Jurafsky. 2016b. Diachronic word embeddings reveal statistical laws of semantic change. *arXiv preprint arXiv:1605.09096*.

Valentin Hofmann, Janet Pierrehumbert, and Hinrich Schütze. 2021a. Dynamic contextualized word embeddings. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 6970–6984.

Valentin Hofmann, Janet Pierrehumbert, and Hinrich Schütze. 2021b. Dynamic contextualized word embeddings. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics*.

Joshua Conrad Jackson, Joseph Watts, Teague R Henry, Johann-Mattis List, Robert Forkel, Peter J Mucha, Simon J Greenhill, Russell D Gray, and Kristen A Lindquist. 2019. Emotion semantics show both cultural variation and universal structure. *Science*, 366(6472):1517–1522.

Ashiqur R KhudaBukhsh, Rupak Sarkar, Mark S Kamlet, and Tom Mitchell. 2021. We don't speak the same language: Interpreting polarization through machine translation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 14893–14901.

Vivek Kulkarni, Rami Al-Rfou, Bryan Perozzi, and Steven Skiena. 2015. Statistically significant detection of linguistic change. In *Proceedings of the 24th international conference on world wide web*, pages 625–635.

Omer Levy, Yoav Goldberg, and Ido Dagan. 2015. Improving distributional similarity with lessons learned from word embeddings. *Transactions of the Association for Computational Linguistics*, 3:211–225.

Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. Efficient estimation of word representations in vector space. *Preprint*, arXiv:1301.3781.

Bill Noble, Asad Sayeed, Raquel Fernández, and Staffan Larsson. 2021. Semantic shift in social networks. In *Proceedings of∗SEM 2021: The Tenth Joint Conference on Lexical and Computational Semantics*, pages 26–37.

Bill Thompson, Seán G Roberts, and Gary Lupyan. 2020. Cultural influences on word meanings revealed through large-scale semantic alignment. *Nature Human Behaviour*, 4(10):1029–1038.